



HH memory and compilation

H200 NVL | optimized HH | B=S=16, T=1600 | float64 | four batches, one worker per method/config

- Leak: $N_\theta = C = N_x/4$ ▲ Leak+K+Na: $N_\theta = 3C = 3N_x/4$ - - RTRL
■ Leak+K: $N_\theta = 2C = N_x/2$ — BPTT



XLA analysis, not process peak GPU memory. Compile includes tracing/lowering; target preparation is excluded.
Lines are guides between measured points; no unmeasured crossover boundary is asserted.